

## Module 10: rstan

To learn about rstan, consider this vignette:

<https://cran.r-project.org/web/packages/rstan/vignettes/rstan.html>

See here for prior choice recommendation:

<https://github.com/stan-dev/stan/wiki/Prior-Choice-Recommendations>

In this notebook, we start simpler, estimating just a mean and variance, followed by the fitting of a hierarchical model. We compare the model used to that used (implicitly) in brm.

### Radon data

Read in the radon data and process (copied from earlier module)

```
# house level data
d <- read.table(url("http://www.stat.columbia.edu/~gelman/arm/examples/radon/srrs2.dat"),
                header=T, sep=",")

# deal with zeros, select what we want, make a fips (county) variable to match on
d <- d %>%
  mutate(activity = ifelse(activity==0, 0.1, activity)) %>%
  mutate(fips = stfips * 1000 + cntyfips) %>%
  dplyr::select(fips, state, county, floor, activity)

# county level data
cty <- read.table(url("http://www.stat.columbia.edu/~gelman/arm/examples/radon/cty.dat"),
                  header = T, sep = ",")

cty <-
  cty %>%
  mutate(fips = 1000 * stfips + cntfips) %>%
  dplyr::select(fips, Uppm) %>%
  rename(ura_county = (Uppm))

dmn <- d %>%
  filter(state=="MN") %>% # Minnesota data only
  dplyr::select(fips, county, floor, activity) %>%
  left_join(cty)

dat <-
  dmn %>%
  mutate(y = log(activity), log_ur = log(ura_county))
head(dat)
```

##	fips	county	floor	activity	ura_county	y	log_ur
## 1	27001 AITKIN		1	2.2	0.502054	0.7884574	-0.6890476

## 2	27001	AITKIN	0	2.2	0.502054	0.7884574	-0.6890476
## 3	27001	AITKIN	0	2.9	0.502054	1.0647107	-0.6890476
## 4	27001	AITKIN	0	1.0	0.502054	0.0000000	-0.6890476
## 5	27003	ANOKA	0	3.1	0.428565	1.1314021	-0.8473129
## 6	27003	ANOKA	0	2.5	0.428565	0.9162907	-0.8473129

## Using rstan to estimate mean and variance

Start simple, let  $y_i = \log(\text{radon})$ , and assume

$$y_i | \mu, \sigma \sim N(\mu, \sigma^2),$$

to estimate  $\mu$  and  $\sigma$ .

We need to define data inputs and a stan model file.

For the stan model file template, go to “File”, “new file”, “stan file”. Note that the default is to estimate  $\mu$  and  $\sigma$  as in our setting! The default template has flat priors, we can add priors. Save it, I called it “module10\_mean\_addpriors.stan”.

Create data inputs:

```
stan_dat <- list(y = dat$y, N = length(dat$y))
```


```
fit <- stan(file = 'module10_stan_mean_addpriors.stan', data = stan_dat)
```

Showing some default outputs, see also the Rstan vignette


```
print(fit)
```

```
## Inference for Stan model: anon_model.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##               mean se_mean   sd    2.5%    25%    50%    75%    97.5% n_eff Rhat
## mu           1.23     0.00 0.03     1.17     1.21     1.23     1.25     1.28  2984    1
## sigma        0.86     0.00 0.02     0.83     0.85     0.86     0.88     0.91  3255    1
## lp__       -329.22     0.03 1.02    -331.81   -329.62   -328.91   -328.49   -328.21  1569    1
##
## Samples were drawn using NUTS(diag_e) at Tue Oct 18 14:09:56 2022.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

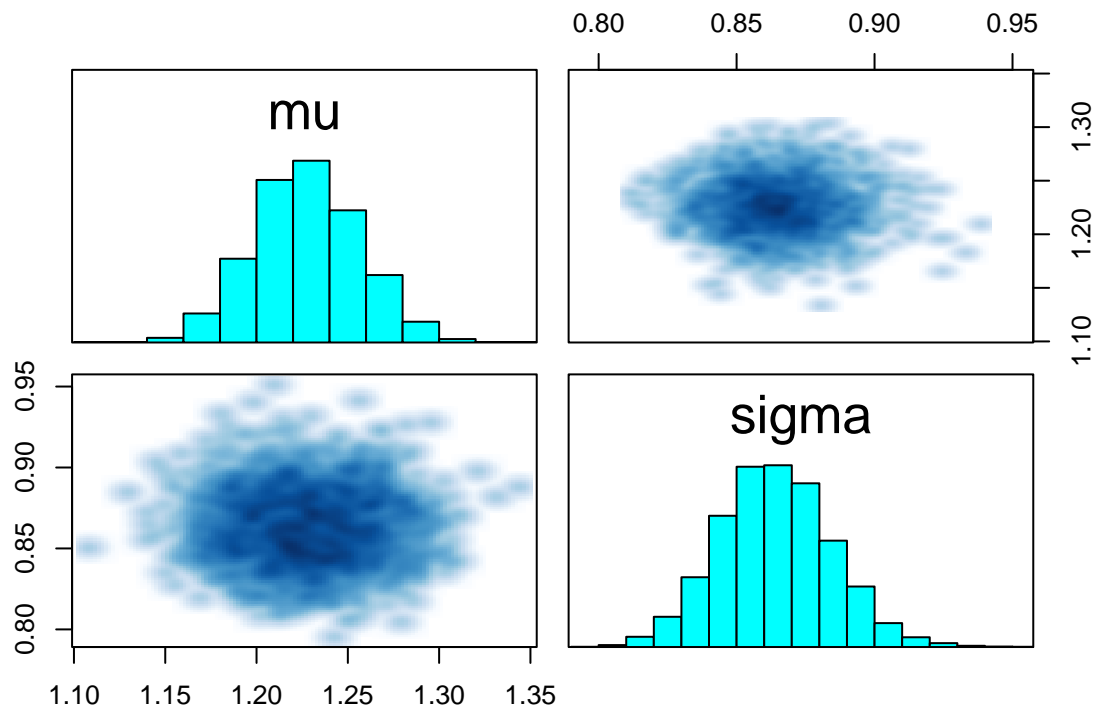
```
plot(fit)
```




```
pairs(fit, pars = c("mu", "sigma"))
```



```
pairs(fit, pars = c("mu", "sigma"))
```



```
traceplot(fit, pars = c("sigma"))
```



## Compare to brm-based model

```
fit_brm <- brm(y ~ 1, data = dat,  
              chains = 4, iter = 1000, warmup = 500, cores = getOption("mc.cores", 4))
```

```
summary(fit_brm)
```

```
## Family: gaussian  
## Links: mu = identity; sigma = identity  
## Formula: y ~ 1  
## Data: dat (Number of observations: 927)  
## Draws: 4 chains, each with iter = 1000; warmup = 500; thin = 1;  
## total post-warmup draws = 2000  
##  
## Population-Level Effects:  
## Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS  
## Intercept 1.23 0.03 1.17 1.29 1.00 1794 1275  
##  
## Family Specific Parameters:  
## Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS  
## sigma 0.86 0.02 0.83 0.91 1.00 1636 1276  
##  
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS  
## and Tail_ESS are effective sample size measures, and Rhat is the potential  
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

```
print(fit)
```

```
## Inference for Stan model: anon_model.  
## 4 chains, each with iter=2000; warmup=1000; thin=1;  
## post-warmup draws per chain=1000, total post-warmup draws=4000.  
##  
## mean se_mean sd 2.5% 25% 50% 75% 97.5% n_eff Rhat  
## mu 1.23 0.00 0.03 1.17 1.21 1.23 1.25 1.28 2984 1  
## sigma 0.86 0.00 0.02 0.83 0.85 0.86 0.88 0.91 3255 1  
## lp__ -329.22 0.03 1.02 -331.81 -329.62 -328.91 -328.49 -328.21 1569 1  
##  
## Samples were drawn using NUTS(diag_e) at Tue Oct 18 14:09:56 2022.  
## For each parameter, n_eff is a crude measure of effective sample size,  
## and Rhat is the potential scale reduction factor on split chains (at  
## convergence, Rhat=1).
```

```
stan_data_from_brm <- standata(fit_brm)  
names(stan_data_from_brm)
```

```
## [1] "N" "Y" "K" "X" "prior_only"
```

```
# N and y (called Y here) are same as in stan_dat  
stan_dat$N
```

```
## [1] 927
```

```
stan_data_from_brm$N
```

```
## [1] 927
```

```
summary(stan_data_from_brm$Y)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -2.3026  0.6419  1.2809  1.2275  1.8245  3.8754
```

```
summary(dat$y)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -2.3026  0.6419  1.2809  1.2275  1.8245  3.8754
```

```
stancode(fit_brm)
```

```
## // generated with brms 2.17.0
## functions {
## }
## data {
##   int<lower=1> N; // total number of observations
##   vector[N] Y; // response variable
##   int prior_only; // should the likelihood be ignored?
## }
## transformed data {
## }
## parameters {
##   real Intercept; // temporary intercept for centered predictors
##   real<lower=0> sigma; // dispersion parameter
## }
## transformed parameters {
##   real lprior = 0; // prior contributions to the log posterior
##   lprior += student_t_lpdf(Intercept | 3, 1.3, 2.5);
##   lprior += student_t_lpdf(sigma | 3, 0, 2.5)
##     - 1 * student_t_lccdf(0 | 3, 0, 2.5);
## }
## model {
##   // likelihood including constants
##   if (!prior_only) {
##     // initialize linear predictor term
##     vector[N] mu = Intercept + rep_vector(0.0, N);
##     target += normal_lpdf(Y | mu, sigma);
##   }
##   // priors including constants
##   target += lprior;
## }
## generated quantities {
##   // actual population-level intercept
##   real b_Intercept = Intercept;
## }
```

See slides re what's going on in this model!

For just the priors

```
get_prior(y~1, data = dat)
```

```
##               prior      class coef group resp dpar nlpar lb ub  source
## student_t(3, 1.3, 2.5) Intercept                                     default
## student_t(3, 0, 2.5)    sigma                                     0    default
```

## Fit a multilevel regression model

Consider the following model:  $y_i | \mu_i, \sigma \sim N(\mu_i, \sigma^2)$ , where

$$\mu_i = \mu_\alpha + \eta_{j[i]} + \beta x_i,$$

where  $j[i]$  refers to county and  $x_i$  to floor indicator.

The stan model for this is given in module10\_hier\_regression.stan.

```
dat <- dat %>%
  mutate(county_id = as.numeric(as_factor(county)))
stan_dat2 <- list(y = dat$y, N = length(dat$y), x = dat$floor, county_id = dat$county_id,
                 J = max(dat$county_id))
```

```
fit2 <- stan(file = 'module10_hier_regression.stan', data = stan_dat2)
```

```
print(fit)
```

```
## Inference for Stan model: anon_model.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##           mean se_mean   sd    2.5%    25%    50%    75%    97.5% n_eff Rhat
## mu         1.23     0.00 0.03     1.17     1.21     1.23     1.25     1.28  2984    1
## sigma      0.86     0.00 0.02     0.83     0.85     0.86     0.88     0.91  3255    1
## lp__    -329.22     0.03 1.02   -331.81 -329.62 -328.91 -328.49 -328.21  1569    1
##
## Samples were drawn using NUTS(diag_e) at Tue Oct 18 14:09:56 2022.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```